Docker Networking A2Z – Masterclass for Developers & DevOps

What is Docker Networking?

- Docker networking allows containers to communicate with:
 - Each other 📞
 - The host machine 🗏
 - The external internet (5)

Docker automatically creates networks and connects containers based on mode.

(2) Key Terms

Term	Meaning Emoj	
Network	Virtual connection b/w containers	
Bridge	Default, isolated internal network	
Host	Shares host's network stack	
None	No network access	0
Overlay Cross-host communication (Swarm)		₩

🔛 Bridge Mode (Default)

Bridge network is like a **private switch** where containers talk to each other.

Z Created Automatically:

docker network ls

Look for: bridge

% How it works:

- Containers get **private IPs** (like 172.17.0.x)
- They can access the internet via NAT
- But cannot be accessed from outside without -p port mapping

Try it:

```
docker run -d --name container1 nginx
docker run -d --name container2 busybox sleep 9999

# Ping container1 from container2 by IP
docker exec -it container2 ping 172.17.0.x
```

X By default, they can't talk by name unless in custom network

E Custom Bridge Network (Recommended)

© Custom networks support container name resolution (DNS)!

Create a custom bridge:

docker network create my-network

Launch containers into it:

```
docker run -d --name app1 --network my-network nginx
docker run -it --name app2 --network my-network busybox sh
```

Now, inside app2:

ping app1

✓ Works! Scontainers can ping by name!

✓ Why Use Custom Bridge?

Feature	Benefit
❷ DNS	Resolve container names
f Isolation	Only containers in same network can talk
S Flexibility	Use multiple networks
Control	Inspect with docker network inspect

2025-06-29 Docker Networking.md



🟠 Host Network Mode

Shares the **host's network stack** directly.



docker run --network host nginx

✓ Pros:

- Super fast no NAT or port mapping
- & Useful for monitoring tools (Prometheus, Grafana)

X Cons:

- \!\text{\Lambda} No isolation
- Eannot run 2 containers on same port!

🤇 None Network Mode

Container has **no networking** at all.

docker run --network none busybox

- Useful for security testing or offline compute jobs

Overlay Network (Advanced – Docker Swarm)

Enables containers on different hosts to communicate

Use Case:

- Docker Swarm
- Distributed Microservices

docker network create --driver overlay my-overlay

Requires Swarm mode:

```
docker swarm init
```

♥ Connect Containers to Multiple Networks

```
docker network create frontend
docker network create backend

docker run -d --name api \
    --network frontend \
    --network-alias api \
    nginx
```

Then attach to another network:

```
docker network connect backend api
```

(2) Inspect a Network

docker network inspect my-network

Shows:

- Container list
- IPs
- Aliases
- Subnets

⊞ CLI Recap

Command

	. u. pose
docker network 1s	List networks
docker network create <name></name>	Create a custom network
docker runnetwork <name></name>	Connect to specific network
docker network inspect <name></name>	Inspect config and members
docker network rm <name></name>	Delete a network
	4.4.40

Purpose

Command	Purpose
docker network connect	Connect running container
docker network disconnect	Disconnect container

& Best Practices for Docker Networking

Practice	Why It's Great
✓ Use custom bridge networks	Enable name resolution + isolation
	Can expose host stack
Use none mode for compute-only jobs	Maximum isolation
Limit network access	Avoid connecting everything together
• Use inspect to debug IPs	Know who talks to whom
Use busybox or alpine to test ping	Lightweight network testing tools

Summary Table

Mode	Isolated?	Can Use DNS?	Host Access?	Notes
bridge	✓ Yes	X No (unless custom)	✓ via -p	Default
custom bridge	✓ Yes	✓ Yes	✓ via -p	Best for local
host	X No	✓ Yes	Direct	No port mapping
none	✓ Yes	X No	X No	For isolation
overlay	✓ Cross-host	✓ Yes	Swarm only	For multi-node

Final Analogy

Bridge network = Private Wi-Fi router ⊕ Custom bridge = Guest Wi-Fi with name tags ⊘ Host mode = Ethernet cable directly into host ■ None mode = Airplane mode ズ Overlay = Corporate VPN connecting multiple offices 🕮 🕮

Docker Networking Overview

Container networking is the foundation of container communication. Every container is equipped with a network interface, IP address, routing table, DNS config, etc.

- ☑ By default, containers can:
 - Make outbound connections (internet)

- Be connected to default or custom networks
- Be isolated or exposed, depending on the configuration

Types of Docker Network Drivers

Network Driver	Description	Use Case
bridge	Default isolated network for single-host	Local development
host	♠ Shares host's network namespace	Low-latency, host-level apps
none	No networking at all	Offline or compute-only tasks
overlay	Transples multi-host (Swarm) communication	Distributed systems
macvlan	■ Assigns MAC + IP from host's network	IoT, legacy systems
ipvlan	Similar to macvlan, IP only	More control over routing
custom plugins	🗱 Extendable third-party solutions	SDN, advanced use

Default Bridge Network vs 😵 User-Defined Bridge Network

Both are based on the bridge driver, but they differ significantly in features & behavior.

Default bridge Network

Created automatically by Docker when installed.

```
docker network ls
# OUTPUT will contain:
# bridge bridge local
```

Example:

```
docker run -d --name app1 nginx
docker run -d --name app2 busybox sleep 999
```

Limitations:

X Limitation	Explanation	
No DNS	Can't resolve container names	
Not isolated	All containers using default bridge are technically on the same LAN	

X Limitation		Explanation	
	Poor discoverability	Need to link manually or use IPs	
1 Less secure		Containers can talk across networks by default if not isolated	

Created with:

docker network create my-custom-net

✓ Advantages:

☑ Feature	Benefit	
Built-in DNS	Containers can resolve each other by name	
solated environment	Containers only talk to others on the same network	
Auto service discovery	Works like microservices	
% Fine-grained control	Inspect, attach, detach easily	
Easier multi-container orchestration	Compose, Swarm, or manually	

& Comparison: Default vs User-Defined Bridge

Feature	Default bridge 🕍	User-Defined Bridge 🕸
DNS support	X No	✓ Yes
Service name resolution	X No	✓ Yes
Isolation	X Less secure	☑ Scoped per network
Compose support	Not recommended	✓ Fully supported
Security	Basic	Scoped & controlled
Container-to-container name access	X Only via IP	✓ Via name (ping app1)
Preferred for production/dev	X No	✓ Yes

Inspecting Networks

docker network inspect my-custom-net

• Output includes:

- Subnet
- Gateway
- Connected containers
- DNS aliases

Example Test

1. Default Bridge (no DNS):

```
docker run -d --name alpha nginx
docker run -it --rm busybox
# ping alpha - X fails
```

2. User-defined Bridge:

```
docker network create testnet
docker run -d --name alpha --network testnet nginx
docker run -it --rm --network testnet busybox
# ping alpha - ☑ works
```

Other Drivers – Quick Overview

Driver	Emoji	Key Use	
host		High-perf apps (no NAT), not isolated	
none	0	Secure offline or processing containers	
overlay	③	Multi-host Swarm networking	
macvlan	r=	Assign physical IPs from host's LAN	
ipvlan	⊛	Fine-grained routing control	
custom plugin		CNI integrations, SDNs (like Calico)	

When to Use What?

Use Case Best Network Driver	
Local dev, isolated apps	😘 User-defined bridge
Multi-container orchestration	
High-speed, low-latency app (e.g. Prometheus)	⚠ Host network

Use Case	Best Network Driver	
No internet access container	None	
Containers across multiple hosts	🕤 Overlay (Swarm mode)	
Assign IPs from LAN for legacy systems	■ Macvlan	

💣 Key Docker Networking Commands Cheat Sheet 🖏

S Command	Description
docker network 1s	📃 List all available networks
docker network create <name></name>	🗱 Create a custom bridge network
docker network rm <name></name>	Remove a network
docker network inspect <name></name>	• View detailed info about a network
docker runnetwork <name> <image/></name>	Create container in a specific network (unnamed)
<pre>docker run -dname <container>network <network> <image/></network></container></pre>	Create named container in a custom network
docker network connect <network> <container></container></network>	Attach an existing container to a network
docker network disconnect <network> <container></container></network>	X Detach a container from a network

Real-World Example: Create and Use a Custom Bridge Network

- # 1 Create a custom bridge network
 docker network create my-bridge
- # 2 Run a container (nginx) attached to that network
 docker run -d --name my-app --network my-bridge nginx
- # 3 Run another container (busybox) in the same network docker run -it --name client --network my-bridge busybox
- # 4 Inside 'client', you can ping 'my-app' by name ping my-app
- Result: client can resolve and communicate with my-app using container name thanks to Docker's internal DNS in custom bridge networks.
- @ Bonus Tip: Disconnect & Reconnect

```
docker network disconnect my-bridge my-app
docker network connect my-bridge my-app
```

Disconnect from network
Reconnect to same or new network

Final Takeaways

- **Default bridge** is basic → no name resolution, low security
- Subser-defined bridge is preferred for real-world apps
- 🖨 Know when to use each driver to optimize performance & security
- Always test communication with tools like ping, curl, netcat

Docker Networking Logic – Container Communication

☑ 1. User-Defined Bridge Network = Container DNS Heaven @

When you create a **user-defined bridge network**, Docker automatically enables an **internal DNS** service.

That means containers can talk to each other by name!

Example:

```
docker network create my-net

docker run -d --name db --network my-net mongo

docker run -d --name web --network my-net node-app
```

& Now, inside web, you can:

```
ping db
```

☑ Boom! It works because Docker auto-resolves db → container's IP inside the same network.

X 2. Default Bridge Network = Only IP-Based Access (12/4)

Containers in the **default bridge** cannot resolve each other by name.

Example:

```
docker run -d --name app1 nginx  # default bridge
docker run -it --name app2 busybox  # default bridge
```

Inside app2:

```
ping app1 # 🗙 FAILS
```

Why? Because **DNS resolution doesn't work** in the default bridge without using the old --link option (now deprecated \circ).

Legacy --link (Avoid Using It)

```
docker run -d --name db mongo
docker run -d --name web --link db node-app
```

Morks — but deprecated & removed in newer Docker versions.

3. Containers in Different Networks = No Communication

Example:

```
docker network create net-a
docker network create net-b

docker run -d --name app1 --network net-a nginx
docker run -d --name app2 --network net-b busybox
```

Inside app2:

```
ping app1 # 🗙 FAILS - isolated networks
```

Metworks are isolated by default. Containers on different bridge networks cannot talk to each other unless you connect one container to both using:

```
docker network connect net-a app2
```

Now, app2 belongs to both networks!

Real-World Analogy

Concept	Analogy
☐ Default Bridge Talking in a crowd with no names, only IPs (♣ 192.168.x.x)	
☼ User-defined Bridge Talking in a chatroom where everyone has a username (@ @we	
ਊ Multi-Network	Having one foot in two rooms 🛆 🛭

Recap – Container Communication Matrix

Scenario	Communicate by name?	Communicate by IP?	Notes
Same user-defined bridge	✓ Yes	✓ Yes	Best practice
Same default bridge	X No	✓ Yes	No name resolution
Different networks	X No	X No	Unless manually connected
Withlink (legacy)	<u></u> Yes	✓ Yes	Deprecated, avoid

✓ Summary

- Substitution User-defined networks allow name-based communication via Docker's built-in DNS.
- **B** Default bridge networks don't support DNS only IPs.
- **(iii)** Each network is isolated containers inside a network can talk, but can't reach containers in other networks unless manually connected.

Overview of Advanced Docker Network Drivers

Driver	Purpose	Host-to- Container	Container-to- Host	Cross-Host Support
overlay	Cross-host container communication (Swarm)	✓ Yes	✓ Yes	✓ Yes
macvlan	Assign real MAC & IP from LAN to container	X No (by default)	✓ Yes	X No
ipvlan	IP-level control without creating MAC addresses	✓ Yes	✓ Yes	X No

1 Overlay Network

What It Is:

Allows **containers on different Docker hosts** to securely communicate as if they were on the same LAN.

Requires Docker Swarm (or other orchestrators). Creates an **encrypted VXLAN tunnel** between hosts.

Real-world Analogy:

Like a VPN that connects branch offices (containers) across cities (hosts).

✓ Use Cases:

- Multi-host microservices
- Docker Swarm services
- HA + distributed architecture

How to Use (with Swarm):

```
# 1. Initialize Swarm
docker swarm init

# 2. Create overlay network
docker network create --driver overlay my-overlay

# 3. Deploy service to use overlay
docker service create \
    --name webapp \
    --replicas 3 \
    --network my-overlay \
    nginx
```

Key Features:

Feature	Benefit
♠ Encrypted traffic	Uses IPSEC tunneling
Auto service discovery	Works across nodes
Built-in load balancing	Container-to-service routing
% Works with Docker Compose (v3+)	Great for multi-host orchestration

2 Macvlan Network

What It Is:

Allows containers to appear as **physical devices on the host's network**, each with their **own IP and MAC**.

- The container **bypasses Docker's NAT**, appearing directly on your LAN.
- Real-world Analogy:
 - Like plugging a new computer (container) directly into your office switch with its own IP.
- ✓ Use Cases:
 - Legacy apps that require static IPs or MACs
 - IoT, embedded, or bare metal simulation
 - When containers must be reachable from the LAN directly


```
# 1. Create macvlan network
docker network create -d macvlan \
    --subnet=192.168.1.0/24 \
    --gateway=192.168.1.1 \
    -o parent=eth0 \
    macvlan-net

# 2. Run a container on that network
docker run -d --name myrouter \
    --network macvlan-net \
    busybox sleep 3600
```

☆ parent=eth0: your physical host's network interface

➤ Limitation ○ No container-to-host		Description
		Can't ping container from host by default
		Be cautious in shared infra
	Requires IP address planning	Must avoid IP conflicts

X Tip to Enable Host ↔ Container Communication (Workaround):

Create a **dummy interface** on the host:

```
ip link add macvlan-shim link eth0 type macvlan mode bridge ip addr add 192.168.1.200/24 dev macvlan-shim ip link set macvlan-shim up
```


What It Is:

Similar to macvlan, but **no extra MACs per container**. All containers **share host's MAC**, just get different IPs.

More compatible with cloud and DHCP setups where duplicate MACs are not allowed.

Analogy:

Multiple workers using one ID card (MAC) but different phone numbers (IP).

✓ Use Cases:

- Performance-sensitive systems
- Cloud infra with MAC restrictions
- · Advanced network routing with minimal overhead

Ø How to Use:

```
docker network create -d ipvlan \
    --subnet=192.168.1.0/24 \
    --gateway=192.168.1.1 \
    -o parent=eth0 \
    ipvlan-net

docker run -it --rm --network ipvlan-net alpine
```

IPvlan supports two modes:

- 12: Same subnet, like macvlan
- 13: Different subnet, routing via host

☑ Benefits:

Feature	Benefit	
@ Efficient	No MAC duplication	

Feature	Benefit	
1 More predictable IP rules	Good for secured environments	
E Custom routing support	Great for advanced network setups	

Advanced Network Drivers Comparison Table

Feature/Driver	overlay 🕤	macvlan 🖷	ipvlan 🎺
Cross-host support	abla	×	×
Requires Swarm?		×	×
Uses physical IP/MAC	×	\checkmark	(IP only)
Host ↔ container communication	$oxed{egin{array}{c} oxed{eta}}$	🗶 (manual fix)	(limited)
Best for	Microservices on multiple hosts	LAN-level communication	Custom IP routing or cloud infra
Security model	Swarm-controlled	Exposes real IP on LAN	More controlled than macvlan
Complexity	Medium	Migh	Medium-High

Security Considerations

Driver	Security Tip
Overlay Isolated per service; enable encryption	
Macvlan	Bypasses Docker's firewall — isolate via VLAN
IPvlan	Good firewall compatibility; still isolate with subnet rules

☑ When to Use What?

Situation	Use Driver
Multi-host Swarm deployments	overlay
☑ Direct LAN access required	macvlan
Controlled IP mapping, no MAC exposure	ipvlan

Summary

- overlay: Best for Swarm, cross-node services, scalable infra
- macvlan: Best for LAN visibility, legacy hardware, IP-bound apps

• ipvlan: Best for performance and controlled environments (e.g., cloud)

O Docker none Network Driver – Ultimate Guide

? What is the none Network?

The none network is a special Docker network driver that **completely disables networking** for a container.

No IP address

No routing

No DNS

No internet access

No communication with host or other containers

No Paddress

No

பி Use Case:

When you want your container to run in complete isolation, especially for:

- CPU-intensive or file-only tasks
- Secure environments with no external communication
- Containers that interact only via volume sharing or IPC
- Avoiding **network-related attacks** (like SSRF, port scanning, etc.)

Real-World Analogy

It's like putting a person in a **soundproof, windowless room** § They can compute, read, or write files – but **cannot talk or hear** the outside world.

How to Use It

```
docker run -it --rm --network none alpine
```

Then inside the container:

```
ping google.com # X Fails
ip addr # Shows no IP
```

☑ The container runs, but it's **completely cut off** from any kind of networking.

Check from the Host

```
docker inspect <container-id> | grep -i "NetworkMode"
# Output: "NetworkMode": "none"
```



Use Case	Why none Works
■ Data processing / math computation	Doesn't need the internet
Security sandboxing	No attack vector via networking
Testing firewall rules	Simulate "no internet" condition
Build-only stages	Avoid leaking credentials over network
Ø DevOps CI/CD	Run isolated build/test tasks with no exposure

Warning

- You cannot ping, curl, apt update, or download anything inside containers using --network none
- Any tools that require internet or inter-container access will fail
- It's **not usable** for most microservices or web APIs

Tip: Combine with Volumes or IPC

If you want to **exchange data** without a network:

```
# Create a shared volume
docker volume create shared-data

# Use it with the isolated container
docker run -it --rm --network none -v shared-data:/data alpine
```

☑ This lets you read/write to shared storage without needing any network access.

Summary Table

Feature	none Driver
IP address	X None
DNS	X None
Host access	X None
Container-to-container	X None
Internet	X None
Use case	Security, sandboxing, isolated compute

final Verdict

If you need absolute network isolation, --network none is your zero-trust go-to option. It's perfect for:

- 🛱 Security-first workloads
- O Disabling all remote calls