# **Docker Orchestration with AWS ECS – Full Guide** (with Node.js Production Testing & Tips)

# What is Docker Orchestration?

**Docker Orchestration** means managing the lifecycle of containers:

- 🕅 Deploying
- ☐ Scaling (up/down)

- A Handling failures

#### Popular Orchestration Tools:

- Docker Swarm
- Kubernetes (K8s) 🛞
- Amazon EKS (for Kubernetes) •

We'll focus on AWS ECS using Fargate (serverless) and EC2 launch type.

# **Step 1: Setting Up AWS Account**

- 1. Go to (3) https://aws.amazon.com/
- 2. Sign up for a free tier account (needs credit/debit card ).
- 3. Enable MFA for security 📆.
- 4. Set region (e.g., us-east-1, ap-south-1) 🕥.
- 5. Create an IAM user with AdministratorAccess if not using root.

## Step 2: Setting up Amazon ECR (Elastic Container Registry)

**ECR** is AWS's private Docker registry.

## **周** Steps:

- 1. AWS Console → Search ECR → Create repository
- 2. Configure:
  - Name: my-app
  - Visibility: Private f
  - Tag immutability: Enabled
- 3. Push Docker Image to ECR:

```
# Authenticate Docker to ECR
aws ecr get-login-password --region us-east-1 | docker login --username AWS --
password-stdin <your-account-id>.dkr.ecr.us-east-1.amazonaws.com
# Build and Tag image
docker build -t my-app .
# Tag with ECR repo URI
docker tag my-app:latest <your-account-id>.dkr.ecr.us-east-1.amazonaws.com/my-
app:latest
# Push image
docker push <your-account-id>.dkr.ecr.us-east-1.amazonaws.com/my-app:latest
```

# **Step 3: Setting up ECS Cluster**

An ECS Cluster is where your containers run.

## **園** Steps:

- 1. AWS Console → ECS → Create Cluster
- 2. Choose:
  - Networking only → Fargate
  - EC2 + Networking → EC2
- 3. Cluster Name: my-app-cluster
- 4. Proceed → ECS will create VPC and subnets 🚷



## Step 4: ECS Task Definition Setup

A **Task Definition** = Docker container blueprint.

### ₩ Key Components:

- Task Role (IAM)
- Docker Image from ECR
- Port mappings (e.g., 80:3000)
- CPU & Memory: 256 CPU, 512 MiB RAM
- Log configuration: AWS CloudWatch
- Environment variables, secrets 🕅
- Health checks: /health or /

### **園** Steps:

1. ECS → Task Definitions → Create new

- 2. Launch Type: Fargate
- 3. Add container:
  - Name: my-app
  - o Image: <ECR Image URL>
  - o Port: 3000
  - Logging: awslogs, with group: /ecs/my-app

# \$\displaystyle \text{Step 5: ECS Service Setup with Load Balancer}\$

#### The **Service** handles:

- Keeping tasks running
- · Restarting failed containers
- Auto-scaling

### **周** Steps:

- 1. ECS → Your Cluster → Create Service
- 2. Launch Type: Fargate
- 3. Choose Task Definition
- 4. Desired tasks: 1 or more
- 5. Attach Load Balancer:
  - ALB → New or existing
  - o Create Target Group → port 3000
  - Health Check path: /health
  - Listener on port 80 → forward to Target Group
- 6. Enable Auto Scaling (optional)

# 

- ☑ Basic Test

  - 2. Visit in browser  $\rightarrow$  You should see your app  $\mathscr{Q}$
  - 3. Confirm task is running: ECS > Cluster > Tasks
  - 4. Logs: CloudWatch > /ecs/my-app
- ✓ Deep Testing (For Node.js Applications)
- ♦ 1. Check container logs:

```
aws logs get-log-events \
    --log-group-name "/ecs/my-app" \
    --log-stream-name "<your-log-stream>"
```

Or via CloudWatch Console.

#### ♦ 2. CURL/HTTP test:

```
curl http://<load-balancer-dns>/health
```

☑ Ensure response is 200 OK. If not, ECS will **kill and restart** your task.

#### **⋄** 3. Test environment variables:

Add this in your Node.js app:

```
console.log('ENV:', process.env.NODE_ENV);
```

Set "NODE\_ENV": "production" in Task Definition.

#### **⋄** 4. Debug failing deployments:

Check:

- Task status (Stopped?)
- View Reason (StoppedReason)
- Logs (CloudWatch)
- Health Check (endpoint must return 2xx)

#### **⋄** 5. Enable ECS Exec:

Run shell commands inside the running container:

```
aws ecs execute-command \
    --cluster my-app-cluster \
    --task <task-id> \
    --container my-app \
    --interactive \
    --command "/bin/sh"
```

Requires enabling ECS Exec & permissions.

#### Checklist:

- Stop ECS Service
- Delete Tasks
- Delete Load Balancer
- Delete Target Groups
- Delete ECR Repository (optional)
- Delete Cluster
- Delete VPC (if created)
- Delete CloudWatch logs

aws ecr delete-repository --repository-name my-app --force

## Manual vs Automatic Orchestration

Feature	Manual	Automatic
Deploy new container $\mathscr{Q}$	CLI or Console	CI/CD + ECS Service Updates
Scale app 🕲	You change task count	Auto Scaling based on CPU/Memory/Requests
Monitor and Heal 🖔	Manual restart	ECS restarts crashed tasks
Load Balancing	Manually configure ELB	ECS auto-registers tasks to Target Groups
Image Updates 😉	Push new tag and update task def	Use CodePipeline / GitHub Actions + Blue/Green

# Real-World Production Tips ( Especially for Node.js)

- 1. Use .env.production with dotenv and pass via Task Definition
- 2. Reverse proxy with NGINX (optional for advanced setups)
- 3. Health Check Endpoint (/health): return 200 OK JSON, no DB calls
- 4. Use pm2 inside container for better process management (optional)
- 5. Avoid console.log in production → Use winston or pino
- 6. Enable Structured Logging → Send logs to CloudWatch
- 7. Monitor memory & CPU metrics via CloudWatch
- 8. Enable ECS Exec to debug running container
- 9. Use HTTPS via ACM with Load Balancer
- 10. Auto-Deploy via GitHub Actions + AWS CLI or CodePipeline
- 11. Set container limits (soft/hard memory) to avoid OOM crashes
- 12. Use Secrets Manager for DB/API credentials 📆
- 13. Test locally with docker run -p 3000:3000 my-app before push
- 14. Set NODE\_ENV=production for optimized performance

#### 15. Use a lightweight base image like node:18-alpine

# Summary Cheatsheet

Step	Task	Tool/Service
1	Create AWS Account	AWS Console
2	Push Image to ECR	ECR, Docker CLI
3	Setup ECS Cluster	ECS
4	Define Task Definition	ECS
5	Create Service + Load Balancer	ECS + ALB
6	Test Your Application	Load Balancer URL
7	Cleanup	Console / CLI